

Joins

What are SQL joins?

→ In SQL, a join is a way to combined data from two or more database tables based on a related column between them.

Joins are used when we want to query information that is distributed across multiple tables in a database, and the information we need is not contained in a single table. By joining tables together, we can create a virtual table that contains all of the information we need for our query.

But why have data in multiple tables?

| userid | name | State | city |
|--------|----------|-------------|-----------|
| 1 | bharat | Gujrat | Ahmedabad |
| 2 | Pearl | Maharashtra | Pune |
| 3 | Jahan | MP | Bhopal |
| 4 | Divsha | Rajasthan | Jaipur |
| 5 | Kasheen | WB | Kolkata |
| 6 | Harzel | Karnatka | Bangalore |
| 7 | Sonakshi | J&K | Kashmir |

| Order id | class id | date |
|----------|----------|------|
| B-25601 | 1 | - |
| B-26011 | 1 | - |
| B-26074 | 1 | - |
| B-25602 | 2 | - |
| B-25114 | 2 | - |
| B-25614 | 2 | - |
| B-22541 | 3 | - |

Common Column

[illegible]

* Why not store two different tables in single large table?

→ Bcz


| user-id | name | state | city |
|---------|------|------------|--------------|
| - | - | <u>Guj</u> | <u>Ahmed</u> |
| - | - | <u>Guj</u> | <u>Ahmed</u> |
| - | - | <u>Guj</u> | <u>Ahmed</u> |
| - | - | <u>Guj</u> | <u>Ahmed</u> |

} — change city, state


we have to change all city and state, if one-two miss than create error in future and take large storage.

Types of JOINS


→ Left JOIN




→ Right Join



→ Inner Join



→ Full Join



→ Cross join

→ Self join

Cross joins → Cartesian Products

In SQL, a cross join (also known as a Cartesian product) is a type of join that returns the Cartesian product of the two tables being joined. In other words, it returns all possible combinations of rows from the two tables.

Cross joins are not commonly used in practice, but they can be useful in certain scenarios, such as generating test data or exploring all possible combinations of items in a product catalogue. However, it's important to be cautious when using cross joins with large tables, as they can generate a very large result set, which can be resource-intensive and slow to process.

Table 1

| col1 | col2 |
|------|------|
| A | 1 |
| B | 2 |
| C | 3 |

3 rows

Table 2

| col3 | col4 |
|------|------|
| X | Y |
| Z | W |

2 rows

Result

| col1 | col2 | col3 | col4 |
|------|------|------|------|
| A | 1 | X | Y |
| A | 1 | Z | W |
| B | 2 | X | Y |
| B | 2 | Z | W |
| C | 3 | X | Y |
| C | 3 | Z | W |

$3 \times 2 = 6$ rows

Inner Joins

In SQL, an inner join is a type of joins operation that combines data from two or more tables based on a specified condition. The inner join returns only the rows from both table that satisfy the specified condition, i.e the matching rows

When you perform an inner join on two tables, the result set will only contain rows where there is a match between the joining columns in both tables. If there is no match, then the row will not be included in the result set.

Table 1

| Employee ID | Name | Department ID | Salary |
|-------------|-------|---------------|--------|
| 1 | John | 1 | 100k |
| 2 | Jane | 2 | 50k |
| 3 | Tim | 3 | 75k |
| 4 | Bob | 1 | 90k |
| 5 | Lisa | 2 | 120k |
| 6 | Mike | 4 | 60k |
| 7 | Sarah | NULL | 80k |





Table 2

| Department ID | Department Name |
|---------------|-----------------|
| 1 | Engineering |
| 2 | Sales |
| 3 | Finance |
| 4 | Marketing |
| 5 | Operations |



| Employee ID | Name | Department ID | Salary | Department |
|-------------|------|---------------|--------|-------------|
| 1 | John | 1 | 100k | Engineering |
| 4 | Bob | 1 | 90k | Engineering |
| 2 | Jane | 2 | 50k | Sales |
| 5 | Lisa | 2 | 120k | Sales |
| 3 | Tim | 3 | 75k | Finance |
| 6 | Mike | 4 | 60k | Marketing |

Left Join

A left join, also known as a left outer join, is a type of SQL join operation that returns all the rows from the left table (also known as the "first" table) and matching rows from the right table (also known as "second" table). If there are no matching rows in the right table, the result will contain NULL values in the columns that come from the right table.

In other words, a left join combines the row from both tables based on a common column, but it also includes all the row from the left table, even if there are no matches in the right table. This is useful when you want to include all the records from the first table, but only some records from the second tables.

| Employee-ID | Name | Department ID | Salary |
|-------------|-------|---------------|--------|
| 1 | John | 1 | 100k |
| 2 | Jane | 2 | 50k |
| 3 | Bob | 3 | 75k |
| 4 | Lisa | 1 | 90k |
| 5 | Mike | 2 | 120k |
| 6 | Tim | 4 | 60k |
| 7 | Sarah | 5 | 80k |
| 8 | Mark | 2 | 95k |





Table 2

| Department ID | Department Name |
|---------------|-----------------|
| 1 | Engineering |
| 2 | Sales |
| 3 | Finance |

| Employee ID | Name | Department ID | Salary | Department ID | Department Name |
|-------------|-------|---------------|--------|---------------|-----------------|
| 1 | John | 1 | 100k | 1 | Engineering |
| 2 | Jane | 2 | 50k | 2 | Sales |
| 3 | Bob | 3 | 75k | 3 | Finance |
| 4 | Lisa | 1 | 90k | 1 | Engineering |
| 5 | Mike | 2 | 120k | 2 | Sales |
| 6 | Tom | 4 | 60k | NULL | NULL |
| 7 | Sarah | 5 | 80k | NULL | NULL |
| 8 | Mark | 2 | 95k | 2 | Sales |

Right Join

A right join, also known as a right outer join, is type of join operations in SQL that returns all the rows from the right table and matching rows from the left tables. If there are no matches in the left table, the result will still contain all the rows from the right table, with Null values for the columns from the left table.




| Employee ID | Name | Department ID | Salary |
|-------------|-------|---------------|--------|
| 1 | John | 1 | 100k |
| 2 | Jane | 2 | 50k |
| 3 | Bob | 3 | 75k |
| 4 | Lisa | 1 | 90k |
| 5 | Mike | 2 | 120k |
| 7 | Sarah | NULL | 80k |
| 8 | Mark | 2 | 95k |
| 8 | | | |

| Department ID | Department Name |
|---------------|-----------------|
| 1 | Engineering |
| 2 | Sales |
| 3 | Financial |
| 4 | Marketing |
| 5 | HR |

| Employee ID | Name | Department ID | Salary | Department Name |
|-------------|------|---------------|--------|-----------------|
| 1 | John | 1 | 100k | Engineering |
| 4 | Lisa | 1 | 90k | Engineering |
| 2 | Jane | 2 | 50k | Sales |
| 5 | Mike | 2 | 120k | Sales |
| 8 | Mark | 2 | 95k | Sales |
| 3 | Bob | 3 | 75k | Finance |
| NULL | NULL | 4 | NULL | Marketing |
| NULL | NULL | 5 | NULL | HR |

Full Outer Join

A full outer join, sometimes called a full join, is a type of join operation in SQL that returns all matching rows from both the left and right tables, as well as any non-matching rows from either table. In other words, a full outer join returns all the rows from both tables and matches rows with common values in the specified column, and fills in NULL values for columns where there is no match.



| emp id | emp name | dept id |
|--------|----------|---------|
| 1 | Alice | 1 |
| 2 | bob | 1 |
| 3 | Charlie | 2 |
| 4 | Dave | null |
| 5 | Eve | 3 |


| dept id | dept name |
|---------|-----------|
| 1 | Sales |
| 2 | Marketing |
| 3 | Finance |
| 4 | IT |
| 5 | HR |

| emp id | emp name | dept id | dept name | dept id |
|--------|----------|---------|-----------|---------|
| 1 | Alice | 1 | Sales | 1 |
| 2 | bob | 1 | Sales | 1 |
| 3 | Charlie | 2 | Marketing | 2 |
| 4 | Dave | null | null | null |
| 5 | Eve | 3 | Finance | 3 |
| null | null | null | IT | 4 |
| null | null | null | HR | 5 |


CROSS JOIN

SELECT * FROM ^{database name} (dvdb1. ^{table name} user) ^{nick name of table} t1

CROSS JOIN ^{database name} dvdb1. ^{table name} groups t2 ^{nick name of table}



INNER TABLE \rightarrow



Show common Table

SELECT * FROM ^{database name} (dvdb1. ^{table name} membership) ^{nick name of table} t1

INNER JOIN ^{table name} dvdb1. user t2


ON t1. user-id = t2. user-id

\hookrightarrow condition

column name
which is common
of both table

* default \rightarrow JOIN is inner join

Left Join




left table and common row will show

SELECT * FROM dvdb1. membership t1

LEFT JOIN dvdb1. user t2


ON t1. user-id = t2. user-id

Right Join →



```
SELECT * FROM dvdbs1. membership t1
RIGHT JOIN dvdbs1. users t2
ON t1. user_id = t2. user_id
```

Full Outer Join →



```
SELECT * FROM dvdbs1. membership t1
FULL OUTER JOIN dvdbs1. user t2
ON t1. user_id = t2. user_id
```


→ This query not working

Join

```
SELECT * FROM dvdbs1. membership t1
LEFT JOIN dvdbs1. user t2
ON t1. user_id = t2. user_id
```

[UNION]

```
SELECT * FROM dvdbs1. membership t1
RIGHT JOIN dvdbs1. user t2
ON t1. user_id = t2. user_id
```



Union
because left join mein common row agya
and Right join mein bhi common row agya.
2 baar common row agya to single bar
karne ke liye UNION.

SQL Set Operation

1. Union: The UNION operator is used to combine the results of two or more SELECT statements into a single result set. The UNION operator removes duplicate rows between the various SELECT statements.

Person 1

| id | Name |
|----|---------|
| 1 | Alice |
| 2 | Bob |
| 3 | Charlie |

Person 2

| id | Name |
|----|---------|
| 3 | Charlie |
| 4 | David |
| 5 | Emily |

Union

| id | name |
|----|---------|
| 1 | Alice |
| 2 | Bob |
| 3 | Charlie |
| 4 | David |
| 5 | Emily |

2. Union All: The union All operator is similar to the UNION operator, but it does not remove duplicate rows from the result set.

| Person 1 | |
|----------|---------|
| Id | name |
| 1 | Alice |
| 2 | Bob |
| 3 | charlie |

| Person 2 | |
|----------|---------|
| Id | Name |
| 3 | charlie |
| 4 | David |
| 5 | Emily |

Union All

| Id | Name |
|----|---------|
| 1 | Alice |
| 2 | Bob |
| 3 | charlie |
| 3 | charlie |
| 4 | David |
| 5 | Emily |

3. INTERSECT: The intersect operator returns only the rows that appear in both result sets of two SELECT statements

| Person 1 | |
|----------|---------|
| Id | name |
| 1 | Alice |
| 2 | Bob |
| 3 | charlie |

| Person 2 | |
|----------|---------|
| Id | Name |
| 3 | charlie |
| 4 | David |
| 5 | Emily |

Intersect

| Id | Name |
|----|---------|
| 3 | charlie |

4. Except: The EXCEPT or MINUS operator return only the distinct row that appear in the first result set but not in the second result set of two SELECT statements.

Person1

| id | name |
|----|---------|
| 1 | Alice |
| 2 | Bob |
| 3 | Charlie |

Person2

| id | name |
|----|---------|
| 3 | Charlie |
| 4 | David |
| 5 | Emily |

Except

| id | name |
|----|--------------------|
| 1 | Alice |
| 2 | Bob Bob |

Union

SELECT * FROM dvdb1. person1

UNION

SELECT * FROM dvdb1. person2

Union All

SELECT * FROM dvdb1. person1

UNION ALL

SELECT * FROM dvdb1. person2

Intersect

SELECT * FROM dvdb1. person1

INTERSECT / EXCEPT

SELECT ~~ALL~~ * FROM dvdb1. person2

Self Joins

A self join is a type of join in which a table is joined with itself. This means that the table is treated as two separate tables, with each row in the table being compared to every other row in the same table.

Self joins are used when you want to compare the values of two different rows within the same table. For example you might use a self join to compare the salaries of two employees who work in the same department, or to find all pairs of customers who have the same billing address.

Table 1

| user-id | name | age | emergency-contact |
|---------|---------|-----|------------------------|
| 1 | Nitish | 34 | 11 → 11 number user-id |
| 2 | Ankit | 32 | 1 → 1 number user-id |
| 3 | Neha | 23 | 1 → " " " " |
| 4 | Radhika | 34 | 3 → 3 num. user-id |
| 8 | Abhinav | 31 | 11 → 11 " " |
| 11 | Rahul | 29 | 8 → 8 " " |

Table 2

| user-id | name | age | emergency-contact |
|---------|---------|-----|-------------------|
| 1 | Nitish | 34 | 11 |
| 2 | Ankit | 32 | 1 |
| 3 | Neha | 23 | 1 |
| 4 | Radhika | 34 | 3 |
| 8 | Abhinav | 31 | 11 |
| 11 | Rahul | 29 | 8 |

Self Join of Table 1 and Table 2

expected output:-

| name | emergency contact for |
|--------|-----------------------|
| Nitish | Rahul → 11 |
| Ankit | Nitish → 1 |
| - | - |
| - | - |
| - | - |

Joining on more than one cols

Student

| Std-id | first-name | last-name | class-id | enroll-year |
|--------|------------|-----------|----------|-------------|
| 1 | John | Smith | 1 | 2021 |
| 2 | Jane | Doe | 2 | 2020 |
| 3 | Bob | Johnson | 1 | 2021 |
| 4 | Sally | Brown | 3 | 2022 |
| 5 | Tom | Williams | 2 | 2022 |
| 6 | Alice | Davis | 4 | 2020 |

Common betⁿ both table but every teacher change. So, we have to match enroll^{year} and class year

Class

| class-id | class-name | teacher | class-year |
|----------|------------|---------------|------------|
| 1 | Math 101 | Mr. Smith | 2021 |
| 2 | English 1 | Mrs. Johnson | 2021 |
| 3 | Science 1 | Dr. Lee | 2022 |
| 4 | History | Mrs. Williams | 2022 |

SELECT * FROM dudb1. students t1

JOIN dudb1. class t2

ON t1. class-id = t2. class-id

AND t1. enrollment-year = t2. class-year

use another joins → Right / Left...

Joining more than 2 tables

SELECT * FROM flipcart. order-detail t1

JOIN flipcart. orders t2

ON t1. order-id = t2. order-id

JOIN flipcart. users t3

ON t2. user-id = t3. user-id

Filtering columns

SELECT

t1. order-id,

t1. amount,

t1. profit,

t3. name

} select ^{ci} specific columns

FROM flipcart. order-detail t1

JOIN flipcart. orders t2

ON t1. order-id = t2. order-id

JOIN flipcart. users t3

ON t2. user-id = t3. user-id